

## Perceptron Algorithm


1. Implement the feature function  $f(x; y)$  with the bag-of-words representations discussed in lecture 2. You can use some preprocessing tricks to reduce the vocabulary size, such as (1) convert all characters into lowercase, and (2) discard low-frequency words or map them into a special token unk.

For this I have taken all the Train set, Dev set and Test set data in variables – trn, trn\_label\_int, dev, dev\_label\_int, tst respectively. I imported stopwords and words from nltk.corpus to filter out stopwords and filter out words not present in words. Also while filtering, I removed special characters and numbers and converted each text to lowercase. I converted the resultant main\_list to a dataframe to calculate the frequency of each word/text and added the frequency as a column to the dataframe by grouping on each unique word. Then finally I removed low frequency words, i.e., frequency < 4, and high frequency words, i.e., frequency > 5000. The resultant vector/feature set was 7715

2. Implement the Perceptron algorithm described in JE section 2.2.1 Algorithm 3.


\_ Plot the accuracy curves on both training and development sets per training epochs, for at least 5 training epochs.

```
: from matplotlib import pyplot as plt
epoch=[1,2,3,4,5]
plt.ylabel('Accuracy Curve')
plt.plot(epoch, train_accuracy)
plt.plot(epoch, dev_accuracy)
plt.show()
```



3. Implement the averaged Perceptron algorithm described in JE section 2.2.2.  
Plot the accuracy curves on both training and development sets per training epochs, for at least 5 training epochs.

```
epoch=[1,2,3,4,5]
plt.ylabel('Accuracy Curve')
plt.plot(epoch, train_avg_accuracy)
plt.plot(epoch, dev_avg_accuracy)
plt.show()
```



## Logistic Regression

1. Use both the LogisticRegression function and the CountVectorizer function with their default settings to train a classifier and report

\_ the size of your feature set - 47963

\_ the classification accuracy on both training and development sets –

- Train\_accuracy \* 100 -> 97.82333333333332
- Dev\_accuracy \* 100 -> 88.0

2. Change the argument ngram range in function CountVectorizer from (1; 1) to (1; 2), then re-train your classi\_er with this large feature set. Report

\_ the size of your feature set - 776704

\_ the classification accuracy on both training and development sets –

- Train\_accuracy \* 100 -> 99.99333333333334
- Dev\_accuracy \* 100 -> 90.44

3. The regularization parameter  $\lambda = 1/c$  in the LogisticRegression function is 1:0. In practice, we need to tune this parameter in order to find the best model. Try different  $\lambda$ 's with the rich feature set built in step 2. For all the  $\lambda$ 's, report the corresponding classification accuracy on both training and development sets.


For  $\lambda$  in [0.961, 0.963, 0.965, 0.967, 0.969, 0.971, 0.973] –

| Train Accuracy    | Dev accuracy |
|-------------------|--------------|
| 99.99333333333334 | 90.45        |
| 99.99333333333334 | 90.45        |
| 99.99333333333334 | 90.46        |
| 99.99333333333334 | 90.46        |
| 99.99333333333334 | 90.45        |
| 99.99333333333334 | 90.45        |
| 99.99333333333334 | 90.45        |


4. Similar to L2 regularization, L1 regularization adds the parameter constraint with its L1 norm. Theoretically, L1 regularization tends to give sparse solutions, which means most of components in  $\theta$  will be 0 or close to 0.

\_In lecture 3, we used contour plots to explain how L2 works. Please use a similar way to explain why L1 regularization prefers sparse solutions.

L1 regularization makes the vector  $x$  smaller(sparse) as most of its elements would be zero/useless and rest elements non-zero and useful. L1 is defined as sum of absolute values of a vectors all elements, i.e.,  $|x| + |y|$ . If we draw all points that have a L1 norm equals to a constant  $c$ , it should be something like this



Only on the axis, the points are sparse, i.e, either  $x$  or  $y$  is zero.



If you notice, the probability of touching a tip is very high. Lp norm when  $0 \leq p < 1$  gives the best result. This reduces the problem of over-fitting.

\_Try different lamda's and report the corresponding classification accuracy on both training and development sets.

For lamda's in [0.963,0.965,0.967,0.968] and penalty='l1', we get –

| Train Accuracy    | Dev accuracy |
|-------------------|--------------|
| 99.29333333333334 | 89.78        |
| 99.29             | 89.81        |
| 99.28333333333333 | 89.80        |
| 99.28             | 89.80        |

5. In lecture 3, we talked different ways to refine the feature set in order to obtain a better classification performance on the development set. Together with the options provided in the LogisticRegression, please try different combinations of these tricks/arguments and \_nd the best model as you can do. Classification accuracy will be an important criterion for evaluating your answers here.

\_ Report the accuracy on both training and development sets with your best model, and explain how you obtain this model.

| Train Accuracy    | Dev accuracy |
|-------------------|--------------|
| 99.91             | 88.41        |
| 99.91             | 88.42        |
| 99.90666666666667 | 88.41        |
| 99.90333333333334 | 88.40        |
| 97.45333333333333 | 88.39        |

1. While vectorizing, via CountVectorizer, converted all text into lowercase, removed words with frequency  $\leq 2$ , ngram\_range=(1,3) and removed stop\_words. This gave me a feature set of 224424.
2. Selected lamda in lamda5=[0.94, 0.95, 0.968, 1, 8.5]
3. For the classifier used  $C=1/\text{lamda}$  and solver='lbfgs' like below -  
`classifier5[i] = LR(C=c_list5[i],solver="lbfgs")`